

Lab 1

Student ID: 1005288

1.1

Part A

```
[09/26/23]seed@VM:~/.../Labsetup$ dockps
e9c439522d9b  hostB-10.9.0.6
0d31473c8a87  hostA-10.9.0.5
5239d41d3812  seed-attacker
[09/26/23]seed@VM:~/.../Labsetup$ █

root@VM:/# ifconfig
br-160cf3e8d0c3: flags=4163<UP,BROADCAST,RUNNING,MULTICAST>  mtu 1500
    inet 10.9.0.1  netmask 255.255.255.0  broadcast 10.9.0.255
    inet6 fe80::42:c7ff:feac:418b  prefixlen 64  scopeid 0x20<link>
ink>

root@e9c439522d9b:/# ping 10.9.0.5 -c 5
PING 10.9.0.5 (10.9.0.5) 56(84) bytes of data.
64 bytes from 10.9.0.5: icmp_seq=1 ttl=64 time=0.062 ms
64 bytes from 10.9.0.5: icmp_seq=2 ttl=64 time=0.059 ms
64 bytes from 10.9.0.5: icmp_seq=3 ttl=64 time=0.065 ms
64 bytes from 10.9.0.5: icmp_seq=4 ttl=64 time=0.059 ms
64 bytes from 10.9.0.5: icmp_seq=5 ttl=64 time=0.083 ms

--- 10.9.0.5 ping statistics ---
5 packets transmitted, 5 received, 0% packet loss, time 4085ms
rtt min/avg/max/mdev = 0.059/0.065/0.083/0.009 ms
```

dockps to find the container IDs and *docksh* to enter the environment.

For this part, I used B to ping A and sniffed the interface br-160cf3e8d0c3 which contains both A and B.

Code was implemented to take in the interface as input as I noticed the interface id changes each time the containers were teared-down and set-up again.

```

###[ Raw ]###
load      = '\x94\x12e\x00\x00\x00\x005\xe2\x0e\x00\x00\x00\x00\x10\x11\x
12\x13\x14\x15\x16\x17\x18\x19\x1a\x1b\x1c\x1d\x1e\x1f !"#%&\'()*+,-./01234567'

###[ Ethernet ]###
dst       = 02:42:0a:09:00:06
src       = 02:42:0a:09:00:05
type      = IPv4
###[ IP ]###
version   = 4
ihl       = 5
tos       = 0x0
len       = 84
id        = 60105
flags     = 
frag      = 0
ttl       = 64
proto     = icmp
chksum    = 0x7bc3
src       = 10.9.0.5
dst       = 10.9.0.6
\options  \
###[ ICMP ]###
type      = echo-reply
code      = 0
chksum    = 0x7e2c
id        = 0x20
seq       = 0x5

```

As shown, I used B(10.9.0.6) to ping A (10.9.0.5), attached is the reply from A to B, and it shows that B successfully pinged A and A replied.

```

[09/26/23]seed@VM:~/.../Lab$ python3 sniffer.py
interface id:br-160cf3e8d0c3
Traceback (most recent call last):
  File "sniffer.py", line 7, in <module>
    pkt = sniff(iface=str(iface),
  File "/usr/local/lib/python3.8/dist-packages/scapy/sendrecv.py", line 1036, in sniff
    sniffer._run(*args, **kwargs)
  File "/usr/local/lib/python3.8/dist-packages/scapy/sendrecv.py", line 906, in _run
    sniff_sockets[L2socket(type=ETH_P_ALL, iface=iface,
  File "/usr/local/lib/python3.8/dist-packages/scapy/arch/linux.py", line 398, in __init__
    self.ins = socket.socket(socket.AF_PACKET, socket.SOCK_RAW, socket.htons(type)) # noqa:
E501
  File "/usr/lib/python3.8/socket.py", line 231, in __init__
    _socket.socket.__init__(self, family, type, proto, fileno)
PermissionError: [Errno 1] Operation not permitted

```

This error was shown after switching to seed and attempting to run the sniffer. This is expected as we require root permissions to sniff packets.

Part B

The code for sniffing ICMP requests is the same as the one in A, hence no modifications were done.

Sniffing tcp connection with specific port number, destination was done by editing filter, filter='tcp and port 23 and src host 10.9.0.6'. The sniffer will sniff tcp connections on port 23, with host B (10.9.0.6)

I used netcat to make A(10.9.0.5) a server listening to tcp requests, and connected B with A on A's IP and port 23 as a client.

Finally, the sniffer was able to pick up the message sent from B to A as shown below.

```
root@a2977f447617:/# nc -v 10.9.0.5 23
Connection to 10.9.0.5 23 port [tcp/telnet] succeeded!
0000 00#00'
Sniffing on Port 23
```

B's message to A

```
len      = 72
id       = 38911
flags    = DF
frag     = 0
ttl      = 64
proto    = tcp
chksum   = 0x8e94
src      = 10.9.0.6
dst      = 10.9.0.5
\options \
###[ TCP ]###
  sport   = 52752
  dport   = telnet
  seq     = 1551095924
  ack     = 590424924
  dataofs = 8
  reserved = 0
  flags   = PA
  window  = 502
  chksum  = 0x1457
  urgptr  = 0
  options = [('NOP', None), ('NOP', None), ('Timestamp', (189025924, 16
74584087))]
###[ Raw ]###
  load    = 'Sniffing on Port 23\n'
```

To simulate a computer on another subnet, I added another host in docker-compose.yml to create another instance of a VM in another subnet with IP 128.230.0.17 as host C. This was done by changing the dockerfile to include a new subnet and a new container for host C.

I used the sniffer to sniff C's traffic as it pings google.com and it was able to capture the packets, which was expected since the attacker and victim had to be in the same subnet / interface. If I were to use an interface of the attacker that did not include C, the sniffer won't sniff any packets.

```

root@VM:/# ifconfig
br-cc70e3a2d9d0: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet 10.9.0.1 netmask 255.255.255.0 broadcast 10.9.0.255
    inet6 fe80::42:b9ff:fe3e:1900 prefixlen 64 scopeid 0x20<link>
    ether 02:42:b9:3e:19:00 txqueuelen 0 (Ethernet)
    RX packets 0 bytes 0 (0.0 B)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 38 bytes 4809 (4.8 KB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

br-dd490ffcc7e2: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet 128.230.0.1 netmask 255.255.0.0 broadcast 128.230.255.255
    inet6 fe80::42:adff:fed7:b77f prefixlen 64 scopeid 0x20<link>
    ether 02:42:ad:d7:b7:7f txqueuelen 0 (Ethernet)
    RX packets 10 bytes 661 (661.0 B)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 48 bytes 5852 (5.8 KB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

```

```

[09/26/23]seed@VM:~/.../A$ dockps
47ec4173acc9 hostC-128.230.0.17
98cdd0e004f8 hostA-10.9.0.5
3ce4f57fb190 seed-attacker
a2977f447617 hostB-10.9.0.6
[09/26/23]seed@VM:~/.../A$ docksh 47
root@47ec4173acc9:/# ifconfig
eth0: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet 128.230.0.17 netmask 255.255.0.0 broadcast 128.230.255.255
    ether 02:42:80:e6:00:11 txqueuelen 0 (Ethernet)
    RX packets 58 bytes 7242 (7.2 KB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 0 bytes 0 (0.0 B)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

lo: flags=73<UP,LOOPBACK,RUNNING> mtu 65536
    inet 127.0.0.1 netmask 255.0.0.0
    loop txqueuelen 1000 (Local Loopback)
    RX packets 0 bytes 0 (0.0 B)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 0 bytes 0 (0.0 B)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

root@47ec4173acc9:/# ping google.com
PING google.com (142.251.175.101) 56(84) bytes of data.
64 bytes from sh-in-f101.1e100.net (142.251.175.101): icmp_seq=1 ttl=101 time=7
.13 ms
64 bytes from sh-in-f101.1e100.net (142.251.175.101): icmp_seq=2 ttl=101 time=6
.69 ms
64 bytes from sh-in-f101.1e100.net (142.251.175.101): icmp_seq=3 ttl=101 time=6
.80 ms
64 bytes from sh-in-f101.1e100.net (142.251.175.101): icmp_seq=4 ttl=101 time=7
.25 ms
64 bytes from sh-in-f101.1e100.net (142.251.175.101): icmp_seq=5 ttl=101 time=1
1.2 ms

```

```

###[ Ethernet ]###
  dst      = 02:42:39:12:43:f2
  src      = 02:42:80:e6:00:11
  type     = IPv4
###[ IP ]###
  version  = 4
  ihl      = 5
  tos      = 0x0
  len      = 84
  id       = 64925
  flags    = DF
  frag     = 0
  ttl      = 64
  proto    = icmp
  chksum   = 0x7db3
  src      = 128.230.0.17
  dst      = 142.251.175.101
  \options \
###[ ICMP ]###
  type     = echo-request
  code     = 0
  chksum   = 0x5036
  id       = 0x1e
  seq      = 0x5
###[ Raw ]###
  load     = '\x92\xd2\x12e\x00\x00\x00\x006\x9c\r\x00\x00\x00\x00\x0
0\x10\x11\x12\x13\x14\x15\x16\x17\x18\x19\x1a\x1b\x1c\x1d\x1e\x1f !"#$%&\'()*+,
-./01234567'

```

1.2

	Time	Source	Destination	Protocol	Length	Info
1	2023-09-26 11:0...	02:42:56:99:b2:fa	Broadcast	ARP	42	Who has 10.9.0.5? Tell 10.9.0.1
2	2023-09-26 11:0...	02:42:0a:09:00:05	02:42:56:99:b2:fa	ARP	42	10.9.0.5 is at 02:42:0a:09:00:05
3	2023-09-26 11:0...	1.2.3.4	10.9.0.5	ICMP	42	Echo (ping) request id=0x0000, seq=0/0,
4	2023-09-26 11:0...	10.9.0.5	1.2.3.4	ICMP	42	Echo (ping) reply id=0x0000, seq=0/0,
5	2023-09-26 11:0...	02:42:0a:09:00:05	02:42:56:99:b2:fa	ARP	42	Who has 10.9.0.1? Tell 10.9.0.5
6	2023-09-26 11:0...	02:42:56:99:b2:fa	02:42:0a:09:00:05	ARP	42	10.9.0.1 is at 02:42:56:99:b2:fa

The original path of the ping was from 10.9.0.1 to 10.9.0.6, but the code spoofed it and the path was changed. The new path is from 1.2.3.4 (non-existent IP) to 10.9.0.5 (another VM on the same network). Wireshark was used to capture the change in paths and a ping request-reply was captured.

1.3

63	2023-09-26	12:3...	216.239.51.16	10.0.2.15	ICMP	70	Time-to-live exceeded (Time to live exceeded in transit)
64	2023-09-26	12:3...	10.0.2.15	74.125.200.100	ICMP	42	Echo (ping) request id=0x0000, seq=0/0, ttl=14 (no response ...)
65	2023-09-26	12:3...	216.239.35.157	10.0.2.15	ICMP	70	Time-to-live exceeded (Time to live exceeded in transit)
66	2023-09-26	12:3...	10.0.2.15	74.125.200.100	ICMP	42	Echo (ping) request id=0x0000, seq=0/0, ttl=15 (no response ...)
67	2023-09-26	12:3...	10.0.2.15	74.125.200.100	ICMP	42	Echo (ping) request id=0x0000, seq=0/0, ttl=16 (no response ...)
68	2023-09-26	12:3...	10.0.2.15	74.125.200.100	ICMP	42	Echo (ping) request id=0x0000, seq=0/0, ttl=17 (no response ...)
69	2023-09-26	12:3...	10.0.2.15	74.125.200.100	ICMP	42	Echo (ping) request id=0x0000, seq=0/0, ttl=18 (no response ...)
70	2023-09-26	12:3...	10.0.2.15	74.125.200.100	ICMP	42	Echo (ping) request id=0x0000, seq=0/0, ttl=19 (no response ...)
71	2023-09-26	12:3...	10.0.2.15	74.125.200.100	ICMP	42	Echo (ping) request id=0x0000, seq=0/0, ttl=20 (no response ...)
72	2023-09-26	12:3...	10.0.2.15	74.125.200.100	ICMP	42	Echo (ping) request id=0x0000, seq=0/0, ttl=21 (no response ...)
73	2023-09-26	12:3...	10.0.2.15	74.125.200.100	ICMP	42	Echo (ping) request id=0x0000, seq=0/0, ttl=22 (no response ...)
74	2023-09-26	12:3...	10.0.2.15	74.125.200.100	ICMP	42	Echo (ping) request id=0x0000, seq=0/0, ttl=23 (no response ...)
75	2023-09-26	12:3...	10.0.2.15	74.125.200.100	ICMP	42	Echo (ping) request id=0x0000, seq=0/0, ttl=24 (reply in 76)
76	2023-09-26	12:3...	74.125.200.100	10.0.2.15	ICMP	60	Echo (ping) reply id=0x0000, seq=0/0, ttl=102 (request ...)
77	2023-09-26	12:3...	10.0.2.15	74.125.200.100	ICMP	42	Echo (ping) request id=0x0000, seq=0/0, ttl=25 (reply in 78)
78	2023-09-26	12:3...	74.125.200.100	10.0.2.15	ICMP	60	Echo (ping) reply id=0x0000, seq=0/0, ttl=102 (request in ...)
79	2023-09-26	12:3...	10.0.2.15	74.125.200.100	ICMP	42	Echo (ping) request id=0x0000, seq=0/0, ttl=26 (reply in 80)
80	2023-09-26	12:3...	74.125.200.100	10.0.2.15	ICMP	60	Echo (ping) reply id=0x0000, seq=0/0, ttl=102 (request in ...)
81	2023-09-26	12:3...	10.0.2.15	74.125.200.100	ICMP	42	Echo (ping) request id=0x0000, seq=0/0, ttl=27 (reply in 82)
82	2023-09-26	12:3...	74.125.200.100	10.0.2.15	ICMP	60	Echo (ping) reply id=0x0000, seq=0/0, ttl=102 (request in ...)

I used a for loop to increment TTL and send the packets to destination (Google's IP address), I used Wireshark to check for ICMP error messages until the TTL is not exceeded. The answer is TTL = 15.

1.4

```
root@VM:/# python3 sniff_and_spoof.py
```

[illegible]

```
seed@VM: ~/.../task1_4
root@d24c7c7fd72b:/# ping 1.2.3.4
PING 1.2.3.4 (1.2.3.4) 56(84) bytes of data.
64 bytes from 1.2.3.4: icmp_seq=1 ttl=64 time=52.6 ms
64 bytes from 1.2.3.4: icmp_seq=2 ttl=64 time=16.9 ms
64 bytes from 1.2.3.4: icmp_seq=3 ttl=64 time=20.6 ms
^C
--- 1.2.3.4 ping statistics ---
3 packets transmitted, 3 received, 0% packet loss, time 2009ms
rtt min/avg/max/mdev = 16.883/30.056/52.638/16.041 ms
root@d24c7c7fd72b:/# ping 8.8.8.8
PING 8.8.8.8 (8.8.8.8) 56(84) bytes of data.
64 bytes from 8.8.8.8: icmp_seq=1 ttl=111 time=6.40 ms
64 bytes from 8.8.8.8: icmp_seq=1 ttl=64 time=21.2 ms (DUP!)
64 bytes from 8.8.8.8: icmp_seq=2 ttl=111 time=7.24 ms
64 bytes from 8.8.8.8: icmp_seq=2 ttl=64 time=18.7 ms (DUP!)
64 bytes from 8.8.8.8: icmp_seq=3 ttl=111 time=6.93 ms
64 bytes from 8.8.8.8: icmp_seq=3 ttl=64 time=19.2 ms (DUP!)
^C
--- 8.8.8.8 ping statistics ---
3 packets transmitted, 3 received, +3 duplicates, 0% packet loss, time 2031ms
rtt min/avg/max/mdev = 6.400/13.268/21.211/6.463 ms
root@d24c7c7fd72b:/# ping 10.9.0.99
PING 10.9.0.99 (10.9.0.99) 56(84) bytes of data.
From 10.9.0.6 icmp_seq=1 Destination Host Unreachable
From 10.9.0.6 icmp_seq=2 Destination Host Unreachable
From 10.9.0.6 icmp_seq=3 Destination Host Unreachable
^C
--- 10.9.0.99 ping statistics ---
5 packets transmitted, 0 received, +3 errors, 100% packet loss, time 4101ms
pipe 4
root@d24c7c7fd72b:/#
```

An inspection of the original sniffed packets shows that icmp.id, icmp.seq and load does not change in the reply and request. As such, we need to

1. Swap dst and src in the IP stack
2. Keep the original load for stack
3. Send to the machine that requested

The results are shown above, the machine got a reply from 1.2.3.4, received duplicates with alerts from 8.8.8.8 and the host is unreachable from 10.9.0.99.

References

<https://dorazaria.github.io/network/packet-sniffing-and-spoofing-lab/>
<https://phoenixnap.com/kb/nc-command#ftoc-heading-2>

